
Benchmark Result

Run ID: BENCH-20260622-0930

Full Benchmark • Model: Qwen2.5-72B-Instruct-AWQ
2026-06-22 11:06

Servers Evaluated: 2
Total Test Cases: 36
Prompts Processed: 4320
Duration: 8100.0s

Contents

1	Executive Summary	3
1.1	Test Configuration	3
1.2	Overall Winner	3
2	Performance Analysis	4
2.1	Time to First Token (TTFT)	4
2.2	Inter-Token Latency (ITL)	4
2.3	Tokens Per Second Per User (TPS)	5
2.4	Request Latency	6
3	Server Comparison	7
3.1	Side-by-Side Comparison	7
3.2	Delta Analysis	7
3.3	Detailed Comparison Table	8
4	Hardware Utilization	10
4.1	Hardware Specifications	10
4.2	Resource Utilization Timeline	10
5	Detailed Breakdown	13
6	Environment & Integrity	14
6.1	Benchmark Environment	14
6.2	Integrity Verification	14

1 Executive Summary

This report presents the benchmark performance analysis for run BENCH-20260622-0930, executed on 2026-06-22 11:06. The benchmark evaluated 2 server(s) using the **Full Benchmark** test suite with the **Qwen2.5-72B-Instruct-AWQ** model.

1.1 Test Configuration

Parameter	Value
Test Suite	Full Benchmark
Model	Qwen2.5-72B-Instruct-AWQ
Servers	192.168.1.10, 192.168.1.20
Tools Used	litellm, ollama_native, oha
Scenarios	short_answer, long_generation, code_generation
Concurrency Levels	1, 8, 16, 32, 64
Total Prompts	4320
Start Time	2026-06-22 08:51:18
End Time	2026-06-22 11:06:18
Duration	8100.0s

1.2 Overall Winner

Based on aggregate performance metrics across all test scenarios, **192.168.1.10** demonstrated superior performance.

2 Performance Analysis

The following charts show key performance metrics across different concurrency levels. Each metric is measured at p50, p95, and p99 percentiles to provide a complete distribution view.

2.1 Time to First Token (TTFT)

TTFT measures the latency from request submission to receiving the first token. Lower values indicate faster initial response.




Figure 1: TTFT by Concurrency Level (p50/p95/p99)

Concurrency	1	8	16	32	64
192.168.1.10					
p50	52.3	98.5	145.2	215.8	385.6
p95	68.5	125.3	185.6	312.5	528.4
p99	85.2	152.8	228.5	425.3	725.8
192.168.1.20					
p50	58.6	112.3	168.5	248.2	435.2
p95	75.2	142.5	215.8	358.6	598.5
p99	95.8	178.5	265.2	485.6	812.3

2.2 Inter-Token Latency (ITL)

ITL measures the average time between consecutive tokens during generation. Lower values indicate smoother streaming.




Figure 2: Inter-Token Latency by Concurrency Level

Concurrency	1	8	16	32	64
192.168.1.10					
p50	18.5	22.3	25.8	31.2	42.5
192.168.1.20					
p50	21.2	25.8	29.5	35.6	48.2

2.3 Tokens Per Second Per User (TPS)

TPS measures the throughput experienced by each individual user. Higher values indicate better throughput efficiency.




Figure 3: Tokens/sec/user by Concurrency Level

Concurrency	1	8	16	32	64
192.168.1.10					
p50	48.5	380.2	720.5	1285.6	2150.8
p95	45.2	352.8	678.5	1185.2	1985.6
p99	42.8	325.6	625.8	1085.3	1825.2
192.168.1.20					
p50	43.2	342.5	648.2	1152.8	1925.6
p95	40.5	315.8	598.5	1052.3	1785.2
p99	38.2	292.5	552.8	958.5	1625.8

2.4 Request Latency

End-to-end request latency from submission to completion.




Figure 4: Request Latency by Concurrency Level (p50/p95/p99)

Concurrency	1	8	16	32	64
192.168.1.10					
p50	125.3	185.6	285.2	452.8	825.6
p95	185.2	298.5	456.8	725.3	1285.6
p99	225.8	385.6	585.2	985.6	1685.2
192.168.1.20					
p50	142.5	215.8	325.6	512.3	945.8
p95	215.6	345.2	528.5	852.6	1456.8
p99	268.5	445.2	685.6	1125.8	1925.6

3 Server Comparison

3.1 Side-by-Side Comparison




Figure 5: Average Performance Metrics by Server

3.2 Delta Analysis

The heatmap below shows the percentage difference between servers for each test configuration. Green indicates improvement, red indicates degradation.

Server Comparison Delta Analysis




Figure 6: Server Comparison Delta Heatmap

3.3 Detailed Comparison Table

Tool	Scenario	C	Server 1				Server 2				Winner
			TTFT	TPS	RPS	P99	TTFT	TPS	RPS	P99	
litellm	short_answer	1	131.07	40.04	7.38	317.49	148.83	35.79	6.69	362.82	server1
litellm	short_answer	8	149.06	320.34	59.04	350.34	174.24	286.30	53.50	407.72	server1
litellm	short_answer	16	169.62	640.68	118.08	387.87	203.28	572.61	107.01	459.04	server1
litellm	short_answer	32	210.74	1281.36	236.16	462.94	261.36	1145.22	214.02	561.69	server1
litellm	short_answer	64	292.98	2562.72	472.32	613.09	377.52	2290.43	428.03	766.97	server1
litellm	long_generation	1	189.01	33.44	2.79	1546.45	216.07	28.93	2.46	1821.41	server1
litellm	long_generation	8	214.95	267.52	22.32	1706.43	252.96	231.47	19.71	2046.84	server1
litellm	long_generation	16	244.60	535.04	44.64	1889.26	295.12	462.94	39.42	2304.49	server1
litellm	long_generation	32	303.89	1070.08	89.28	2254.93	379.44	925.89	78.85	2819.78	server1
litellm	long_generation	64	422.48	2140.16	178.56	2986.26	548.08	1851.78	157.70	3850.36	server1
litellm	code_generation	1	155.14	36.95	4.86	835.85	182.96	31.83	4.31	962.72	server1
litellm	code_generation	8	176.44	295.64	38.88	922.32	214.20	254.66	34.50	1081.88	server1
litellm	code_generation	16	200.77	591.28	77.76	1021.14	249.90	509.31	68.99	1218.06	server1
litellm	code_generation	32	249.44	1182.56	155.52	1218.78	321.30	1018.62	137.98	1490.42	server1
litellm	code_generation	64	346.79	2365.12	311.04	1614.06	464.10	2037.25	275.97	2035.15	server1
ollama_native	short_answer	1	114.55	43.03	8.19	289.88	135.91	37.90	7.30	331.66	server1
ollama_native	short_answer	8	130.27	344.28	65.52	319.87	159.12	303.23	58.43	372.72	server1
ollama_native	short_answer	16	148.24	688.56	131.04	354.14	185.64	606.46	116.86	419.63	server1
ollama_native	short_answer	32	184.17	1377.12	262.08	422.69	238.68	1212.93	233.73	513.46	server1
ollama_native	short_answer	64	256.04	2754.24	524.16	559.78	344.76	2425.86	467.46	701.12	server1
ollama_native	long_generation	1	171.87	35.91	3.06	1405.98	200.08	31.19	2.64	1655.06	server1
ollama_native	long_generation	8	195.46	287.28	24.48	1551.42	234.24	249.50	21.12	1859.92	server1
ollama_native	long_generation	16	222.42	574.56	48.96	1717.65	273.28	499.01	42.24	2094.03	server1
ollama_native	long_generation	32	276.34	1149.12	97.92	2050.10	351.36	998.02	84.48	2562.26	server1
ollama_native	long_generation	64	384.18	2298.24	195.84	2714.99	507.52	1996.03	168.96	3498.72	server1
ollama_native	code_generation	1	141.47	38.48	5.22	767.64	166.36	33.86	4.58	894.31	server1
ollama_native	code_generation	8	160.89	307.80	41.76	847.06	194.76	270.85	36.61	1005.00	server1
ollama_native	code_generation	16	183.08	615.60	83.52	937.81	227.22	541.70	73.22	1131.51	server1
ollama_native	code_generation	32	227.47	1231.20	167.04	1119.32	292.14	1083.39	146.43	1384.52	server1
ollama_native	code_generation	64	316.24	2462.40	334.08	1482.35	421.98	2166.78	292.86	1890.53	server1
oha	short_answer	1	121.28	41.89	7.92	302.67	141.86	37.08	7.04	348.26	server1
oha	short_answer	8	137.92	335.16	63.36	333.98	166.08	296.61	56.32	391.36	server1
oha	short_answer	16	156.95	670.32	126.72	369.77	193.76	593.22	112.64	440.62	server1

Tool	Scenario	C	Server 1				Server 2				Winner
			TTFT	TPS	RPS	P99	TTFT	TPS	RPS	P99	
oha	short_answer	32	195.00	1340.64	253.44	441.34	249.12	1186.43	225.28	539.15	server1
oha	short_answer	64	271.09	2681.28	506.88	584.47	359.84	2372.86	450.56	736.20	server1
oha	long_generation	1	179.32	34.67	2.88	1478.65	207.56	30.18	2.55	1729.07	server1
oha	long_generation	8	203.93	277.40	23.04	1631.62	243.00	241.41	20.42	1943.08	server1
oha	long_generation	16	232.06	554.80	46.08	1806.43	283.50	482.82	40.83	2187.67	server1
oha	long_generation	32	288.31	1109.60	92.16	2156.06	364.50	965.63	81.66	2676.84	server1
oha	long_generation	64	400.82	2219.20	184.32	2855.33	526.50	1931.26	163.33	3655.17	server1
oha	code_generation	1	148.10	37.72	5.04	801.24	175.07	33.03	4.40	928.72	server1
oha	code_generation	8	168.43	301.72	40.32	884.13	204.96	264.22	35.20	1043.67	server1
oha	code_generation	16	191.66	603.44	80.64	978.86	239.12	528.45	70.40	1175.04	server1
oha	code_generation	32	238.13	1206.88	161.28	1168.31	307.44	1056.90	140.80	1437.78	server1
oha	code_generation	64	331.06	2413.76	322.56	1547.22	444.08	2113.79	281.60	1963.27	server1

4 Hardware Utilization

4.1 Hardware Specifications

Metric	192.168.1.10	192.168.1.20
GPU		
GPU Model	NVIDIA RTX 4090 24GB	NVIDIA RTX 3090 24GB
GPU Util (avg)	87.3	92.1
GPU Util (max)	99.8	100.0
VRAM Used (avg)	21.50 GB	22.80 GB
VRAM Total	24.00 GB	24.00 GB
Power (avg)	285.6 W	310.2 W
Power (max)	350.0 W	350.0 W
Temp (avg)	72.5	78.3
Temp (max)	83.0	88.0
CPU / RAM		
CPU (avg)	35.2	42.8
CPU (max)	68.5	75.2
RAM Used (avg)	28.40 GB	32.10 GB
RAM Total	64.00 GB	64.00 GB
Disk / Network		
Disk Read (avg)	125.30 MB/s	98.70 MB/s
Disk Write (avg)	42.80 MB/s	35.20 MB/s

4.2 Resource Utilization Timeline




Figure 7: GPU Utilization Over Time




Figure 8: CPU Utilization Over Time




Figure 9: RAM Usage Over Time




Figure 10: VRAM Usage Over Time

5 Detailed Breakdown

Performance metrics grouped by tool and test scenario.

Tool	Scenario	192.168.1.10							192.168.1.20						
		TPS	TTFT	TPOT	RPS	P99	Err%		TPS	TTFT	TPOT	RPS	P99	Err%	
litellm	short_answer	42.15	128.50	23.40	8.20	312.80	0%		38.90	145.20	25.10	7.60	356.40	0%	
litellm	long_generation	35.20	185.30	28.70	3.10	1523.60	1%		31.45	210.80	31.20	2.80	1789.20	2%	
litellm	code_generation	38.90	152.10	25.80	5.40	823.50	0%		34.60	178.50	28.90	4.90	945.70	1%	
ollama_native	short_answer	45.30	112.30	21.80	9.10	285.60	0%		41.20	132.60	23.50	8.30	325.80	0%	
ollama_native	long_generation	37.80	168.50	26.40	3.40	1385.20	1%		33.90	195.20	29.50	3.00	1625.80	2%	
ollama_native	code_generation	40.50	138.70	24.60	5.80	756.30	0%		36.80	162.30	27.10	5.20	878.50	0%	
oha	short_answer	44.10	118.90	22.50	8.80	298.20	0%		40.30	138.40	24.20	8.00	342.10	0%	
oha	long_generation	36.50	175.80	27.30	3.20	1456.80	1%		32.80	202.50	30.50	2.90	1698.50	2%	
oha	code_generation	39.70	145.20	25.20	5.60	789.40	0%		35.90	170.80	27.80	5.00	912.30	1%	

6 Environment & Integrity

6.1 Benchmark Environment

- **Platform:** aiDaptive Benchmark Suite v2.1
- **Execution Time:** 2026-06-22 08:51:18 – 2026-06-22 11:06:18
- **Prompt Set:** Vietnamese QA v2
- **Concurrency Levels:** 1, 8, 16, 32, 64

6.2 Integrity Verification

- **Raw Tool Output:** Captured & verified for each test case
- **Per-Prompt Logging:** Enabled (TTFT, TPS, TPOT per request)
- **Audit Package:** Available as ZIP with SHA-256 signed MANIFEST
- **Cryptographic Hash:** SHA-256 validation enabled

Note: This report is auto-generated from raw benchmark tool outputs. To verify the cryptographic integrity, download the Evidence Package (ZIP) which includes a MANIFEST.json with per-file SHA-256 hashes.

Full Benchmark

— End of Report —